

About Project: IMDB Movie Analysis

Objective:

As a data analyst intern at IMDB, you have been tasked with exploring and analyzing the IMDB Movies dataset. Your goal is to answer specific business questions, gain insights into movie trends, and deliver actionable recommendations. Using Python and libraries such as Pandas, NumPy, Seaborn, and Matplotlib, perform analysis to help IMDB better understand genre popularity, rating trends, and factors influencing movie success.

Tools and Libraries Used

- **Python**
- **Pandas**: Data manipulation and analysis
- **NumPy**: Numerical computations
- **Matplotlib**: Data visualization
- **Seaborn**: Advanced visualization

About Company

IMDb (Internet Movie Database) is a comprehensive online database of information about films, television shows, video games, and online streaming content. It includes details such as cast and crew, plot summaries, user reviews, trivia, and ratings. Established in 1990, IMDb has become one of the most popular platforms for movie enthusiasts and industry professionals alike. It features user-generated content, professional critiques, and a proprietary rating system based on user votes. Owned by Amazon since 1998, IMDb also offers a subscription service, IMDbPro, providing industry-focused features like contact information and production updates.

Dataset Overview

The dataset includes the following columns:

- **names**: Movie titles
- **date_x**: Release dates
- **score**: IMDB ratings

- **genre**: Genres
- **overview**: Movie summaries
- **crew**: Cast and crew information
- **orig_title**: Original titles
- **status**: Release status (e.g., released, post-production)
- **orig_lang**: Original language
- **budget_x**: Production budgets
- **revenue**: Box office revenues
- **country**: Production country

Loading the dataset and Perform initial setup

Task: Load the dataset and perform initial setup

```
# Importing necessary libraries for the project
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Loading the dataset
```

```
Data = "imdb_movies.csv"
df = pd.read_csv(Data)
```

```
# Display the top 5 rows of datasets
```

```
df.head()
```

	names	date_x	score	\
0	Creed III	03/02/2023	73.0	
1	Avatar: The Way of Water	12/15/2022	78.0	
2	The Super Mario Bros. Movie	04/05/2023	76.0	
3	Mummies	01/05/2023	70.0	
4	Supercell	03/17/2023	61.0	

```

                                genre \
0                                Drama, Action
1                                Science Fiction, Adventure, Action
2                                Animation, Adventure, Family, Fantasy, Comedy
3                                Animation, Comedy, Family, Adventure, Fantasy
4                                Action

                                overview \
0                                After dominating the boxing world, Adonis Cree...
1                                Set more than a decade after the events of the...
2                                While working underground to fix a water main,...
3                                Through a series of unfortunate events, three ...
4                                Good-hearted teenager William always lived in ...

                                crew \
0                                Michael B. Jordan, Adonis Creed, Tessa Thompso...
1                                Sam Worthington, Jake Sully, Zoe Saldaña, Neyt...
2                                Chris Pratt, Mario (voice), Anya Taylor-Joy, P...
3                                Óscar Barberán, Thut (voice), Ana Esther Albor...
4                                Skeet Ulrich, Roy Cameron, Anne Heche, Dr Quin...

                                orig_title    status    orig_lang
budget_x \
0                                Creed III    Released    English
75000000.0
1                                Avatar: The Way of Water    Released    English
460000000.0
2                                The Super Mario Bros. Movie    Released    English
100000000.0
3                                Momias    Released    Spanish, Castilian
12300000.0
4                                Supercell    Released    English
77000000.0

                                revenue country
0    2.716167e+08    AU
1    2.316795e+09    AU
2    7.244590e+08    AU
3    3.420000e+07    AU
4    3.409420e+08    US

```

Data Overview and Basic Exploration

Task: Explore the structure and composition of the dataset

```

# Display the shape and first few rows of the dataset
print("Shape of the dataset:", df.shape)

df.info()

```

```

Shape of the dataset: (10178, 12)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10178 entries, 0 to 10177
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   names           10178 non-null  object
1   date_x          10178 non-null  object
2   score           10178 non-null  float64
3   genre           10093 non-null  object
4   overview        10178 non-null  object
5   crew            10122 non-null  object
6   orig_title      10178 non-null  object
7   status          10178 non-null  object
8   orig_lang       10178 non-null  object
9   budget_x        10178 non-null  float64
10  revenue         10178 non-null  float64
11  country         10178 non-null  object
dtypes: float64(3), object(9)
memory usage: 954.3+ KB

```

- **Rows:** Each row represents a unique movie and contains details like its title, genre, release date, rating, and other attributes.
- **Columns:** Each column represents a feature or attribute of the movies, such as `genre`, `budget_x`, `revenue`, etc.

```

#The issue is with the date_x column as its dtype is given as object
#Convert data type of date_x into datetime
df["date_x"] = pd.to_datetime(df["date_x"])

```

```

# Checking the datatype again
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10178 entries, 0 to 10177
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   names           10178 non-null  object
1   date_x          10178 non-null  datetime64[ns]
2   score           10178 non-null  float64
3   genre           10093 non-null  object
4   overview        10178 non-null  object
5   crew            10122 non-null  object
6   orig_title      10178 non-null  object
7   status          10178 non-null  object
8   orig_lang       10178 non-null  object
9   budget_x        10178 non-null  float64
10  revenue         10178 non-null  float64

```

```

11 country      10178 non-null object
dtypes: datetime64[ns](1), float64(3), object(8)
memory usage: 954.3+ KB

```

```

# Summary statistics for numerical columns
df.describe()

```

	date_x	score	budget_x
revenue			
count	10178	10178.000000	1.017800e+04
1.017800e+04			
mean	2008-06-15 06:16:37.445470720	63.497052	6.488238e+07
2.531401e+08			
min	1903-05-15 00:00:00	0.000000	1.000000e+00
0.000000e+00			
25%	2001-12-25 06:00:00	59.000000	1.500000e+07
2.858898e+07			
50%	2013-05-09 00:00:00	65.000000	5.000000e+07
1.529349e+08			
75%	2019-10-17 00:00:00	71.000000	1.050000e+08
4.178021e+08			
max	2023-12-31 00:00:00	100.000000	4.600000e+08
2.923706e+09			
std	NaN	13.537012	5.707565e+07
2.777880e+08			

```

# Checking for missing values and their counts
df.isnull().sum()

```

```

names      0
date_x      0
score      0
genre      85
overview    0
crew       56
orig_title  0
status      0
orig_lang   0
budget_x    0
revenue     0
country     0
dtype: int64

```

Data Cleaning

Task: Address missing values, data types, and outliers.

```

#Missing Values
#Genre column have 85 missing values which contains Genres of movies

```

```
#Filling Genre with "Unavaiable" as its an object type
df["genre"]=df["genre"].fillna("unavaiable")

#Crew column have 56 missing values which contains Cast and crew
information
#Filling it same with "Unavaiable" as its also an object type
df["crew"]=df["crew"].fillna("unavaiable")
```

Univariate Analysis: Explore each column individually

Task: Perform univariate analysis on numerical and categorical variables


```
# Analyze the distribution of IMDB ratings (`score`) using a histogram
and describe its shape
```

```
# Plotting the distribution of IMDB scores
```

```
plt.figure(figsize=(12, 4))
sns.histplot(df["score"], kde=True, bins=20, color="green")
plt.title("Distribution of IMDB Scores")
plt.xlabel("IMDB Score")
plt.ylabel("Frequency")
plt.show()
```

```
# Description of the shape
```

```
print("The distribution of IMDB scores shows that:")
print("- It appears to be approximately normal but slightly left-
skewed (tail on the lower end).")
print("- Most scores are concentrated between 50 and 80, with a
gradual decline outside this range.")
print("- There are very few movies with extremely low or perfect
scores (0 or 100).")
```



The distribution of IMDB scores shows that:

- It appears to be approximately normal but slightly left-skewed (tail

on the lower end).

- Most scores are concentrated between 50 and 80, with a gradual decline outside this range.
- There are very few movies with extremely low or perfect scores (0 or 100).

What are the most common genres in the dataset? Use a bar chart to show their distribution.

Group by genre and count the number of movies


```
gb = df.groupby("genre").agg({"names": "count"})  
gb = gb.sort_values(by="names", ascending=False)  
gb = gb.head(20)
```

Plotting the bar chart

```
plt.figure(figsize=(12, 4))  
sns.barplot(x=gb.index, y=gb["names"], data=gb, hue=gb.index)  
plt.title("Most Common Genres")  
plt.xlabel("Genre")  
plt.ylabel("Count")  
plt.xticks(rotation=90)  
plt.show()
```

Print the most common genre

```
most_common_genre = gb.index[0]  
print(f"The most common genre in the dataset is: {most_common_genre}")
```




The most common genre in the dataset is: Drama

Bivariate Analysis: Explore relationships between two variables

Task: Use scatter plots, box plots, and correlation analysis

```
# Explore the relationship between `budget_x` and `revenue` using a scatter plot.
# Plotting the Scatter plot
plt.figure(figsize=(12,3))
sns.scatterplot(x="budget_x", y = "revenue",data = df, hue = "budget_x", palette = "viridis")
plt.title("Budget vs Revenue")
plt.xlabel("Budget")
plt.ylabel("Revenue")
plt.show()

# Description observed trend
print("The relationship of budget and revenue shows that:")
print("- There is a positive correlation between budget and revenue. Movies with higher budgets tend to generate higher revenues.")
print("- Most movies are clustered around mid-range budgets (e.g., 80 million).")
print("- Outliers show both successes (high revenue, high budget) and failures (high budget, low revenue).")
```



The relationship of budget and revenue shows that:

- There is a positive correlation between budget and revenue. Movies with higher budgets tend to generate higher revenues.
- Most movies are clustered around mid-range budgets (e.g., 80 million).
- Outliers show both successes (high revenue, high budget) and failures (high budget, low revenue).

```
# Compare IMDB ratings (`score`) across Country using a boxplot.
```

```
# Boxplot for IMDB score across countries
```

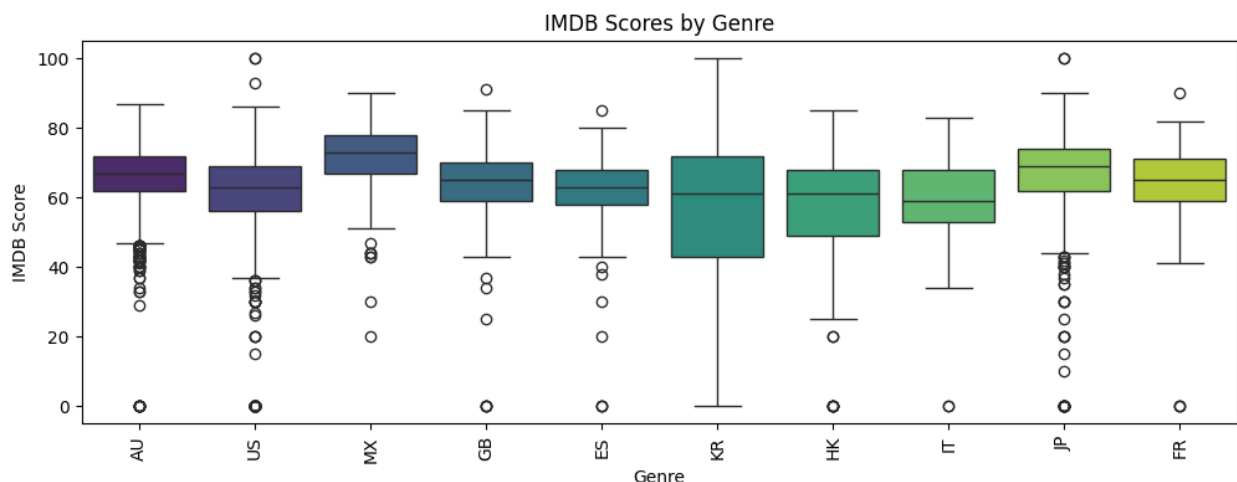


```

top_countries = df["country"].value_counts().head(10).index # Top 10
countries by movie count
plt.figure(figsize=(12, 4))
sns.boxplot(data=df[df["country"].isin(top_countries)], x="country",
y="score", hue="country", palette="viridis", legend=False)
plt.title("IMDB Scores by Genre")
plt.xlabel("Genre")
plt.ylabel("IMDB Score")
plt.xticks(rotation=90)
plt.show()

# Description
print("The IMDB score across contries:")
print("- The distribution of ratings varies significantly across
countries.")
print("- Some countries exhibit tighter distributions, while others
have a wider range of scores.")
print("-The median rating differs by country, providing insight into
regional preferences and quality of movies.")

```



The IMDB score across contries:

- The distribution of ratings varies significantly across countries.
- Some countries exhibit tighter distributions, while others have a wider range of scores.
- The median rating differs by country, providing insight into regional preferences and quality of movies.

#Is there a correlation between the number of votes a movie received and its rating? Create a scatter plot and calculate the correlation coefficient.

Create a dataframe to store score, budget and revenue and find it's correlation

```

df1 = df[["score", "budget_x", "revenue"]]
df1.columns = ["Score", "Budget", "Revenue"]

```

```
corr = df1.corr()
corr
```

	Score	Budget	Revenue
Score	1.000000	-0.23547	0.096533
Budget	-0.235470	1.00000	0.673830
Revenue	0.096533	0.67383	1.000000

```
#correlation between `budget_x`, 'revenue' and `score`.
```

```
correlation = df[["score","budget_x","revenue"]]
```

```
data = correlation.corr()
```

```
sns.heatmap(data, annot = True)
```

```
plt.show()
```

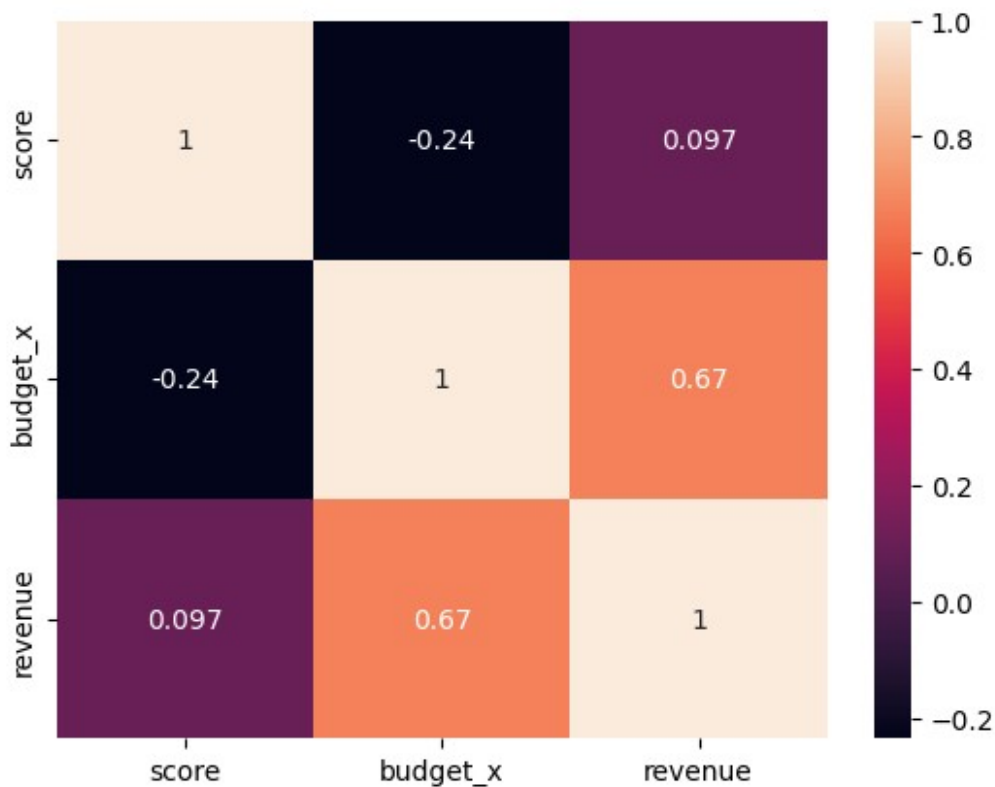
```
# Description
```

```
print("The Correlation between budget, score and revenue:")
```

```
print("- Score and Budget: A weak negative correlation of -0.235,  
meaning as the budget increases, the score slightly decreases..")
```

```
print("- Score and Revenue: A very weak positive correlation of 0.097,  
suggesting little to no relationship between score and revenue.")
```

```
print("-Budget and Revenue: A moderate positive correlation of 0.674,  
indicating that as the budget increases, revenue tends to increase as  
well")
```



The Correlation between budget, score and revenue:

- Score and Budget: A weak negative correlation of -0.235, meaning as the budget increases, the score slightly decreases..
- Score and Revenue: A very weak positive correlation of 0.097, suggesting little to no relationship between score and revenue.
- Budget and Revenue: A moderate positive correlation of 0.674, indicating that as the budget increases, revenue tends to increase as well

Genre-Specific Analysis

Task: Delve deeper into the genre of movies

```
# Which genre has the highest average rating? Calculate the average
rating for each genre and plot the results.
```


```
# Calculate the average rating for each genre
avg_rating_by_genre= df.groupby("genre")
["score"].mean().sort_values(ascending=False)
avg_rating_by_genre = avg_rating_by_genre.head(10)
print("\nAverage Ratings by Genre:")
print(avg_rating_by_genre)
```

Average Ratings by Genre:

```
genre
Adventure, Animation, Comedy, Fantasy, Mystery
85.000000
Fantasy, Drama, Crime
85.000000
Romance, Animation, Drama
85.000000
Animation, Comedy, Romance
84.666667
Animation, Thriller
83.000000
TV Movie, Animation, Science
Fiction, Action, Adventure, Comedy, Drama, Fantasy, Music      83.000000
Animation, Action, Adventure, Fantasy, Thriller
83.000000
Animation, Action, Science Fiction, Drama
83.000000
Family, Animation, Drama
83.000000
Adventure, Fantasy, Animation
83.000000
Name: score, dtype: float64
```

```
# Plot the graph of average score across genres
plt.figure(figsize=(12, 4))
sns.barplot(x=avg_rating_by_genre.index, y=avg_rating_by_genre.values,
hue=avg_rating_by_genre.index, palette='viridis', legend=False)
plt.title("Average IMDB Ratings by Genre")
plt.xlabel("Genre")
plt.ylabel("Average Rating")
plt.xticks(rotation=90)
plt.show()

# Description
print("The Average rating for each genre shows that:")
print("- It appears that (Fantasy, Drama, Crime)(Adventure, Animation,
Comedy, Fantasy, Mystery)(Romance, Animation, Drama) have highest
score.")
```



The Average rating for each genre shows that:

- It appears that (Fantasy, Drama, Crime)(Adventure, Animation, Comedy, Fantasy, Mystery)(Romance, Animation, Drama) have highest score.

How does the popularity of genres vary over time? Plot the number of movies released per genre each year.

```
# Convert 'date_x' to datetime format if it's not already
df['date_x'] = pd.to_datetime(df['date_x'], errors='coerce')
```

```
# Extract the year from the 'date_x' column
```

```

df['year'] = df['date_x'].dt.year

# Check if the 'year' column is created correctly
print(df[['date_x', 'year']].head())


    date_x  year
0 2023-03-02 2023
1 2022-12-15 2022
2 2023-04-05 2023
3 2023-01-05 2023
4 2023-03-17 2023

# Group by 'genre' and 'year', then count movies
movies_per_genre = df.groupby(['year',
'genre']).size().reset_index(name='movie_count')
top_genres = movies_per_genre.groupby('genre')
['movie_count'].sum().nlargest(10).index

# Filter the movies data to include only these top genres
filtered_movies =
movies_per_genre[movies_per_genre['genre'].isin(top_genres)]

# Plotting Line
plt.figure(figsize=(12,4))
sns.lineplot(data=filtered_movies, x='year', y='movie_count',
hue='genre', marker='o')
plt.title("Number of Movies Released Per Genre Each Year")
plt.xlabel("Year")
plt.ylabel("Number of Movies Released")
plt.xticks(rotation=45)
plt.show()

```



```

# Compare budgets and revenues for specific genres.

```

```
# Filter data for specific genres
selected_genres = ['Action', 'Comedy', 'Drama']
df_selected_genres = df[df['genre'].isin(selected_genres)]

# Check the first few rows of the filtered data
print(df_selected_genres[['genre', 'budget_x', 'revenue']].head())
```

	genre	budget_x	revenue
4	Action	77000000.0	340941958.6
22	Drama	25000000.0	622313635.0
42	Action	1000000.0	22790427.2
44	Drama	3000000.0	36752010.0
54	Comedy	51880004.0	9263009.0

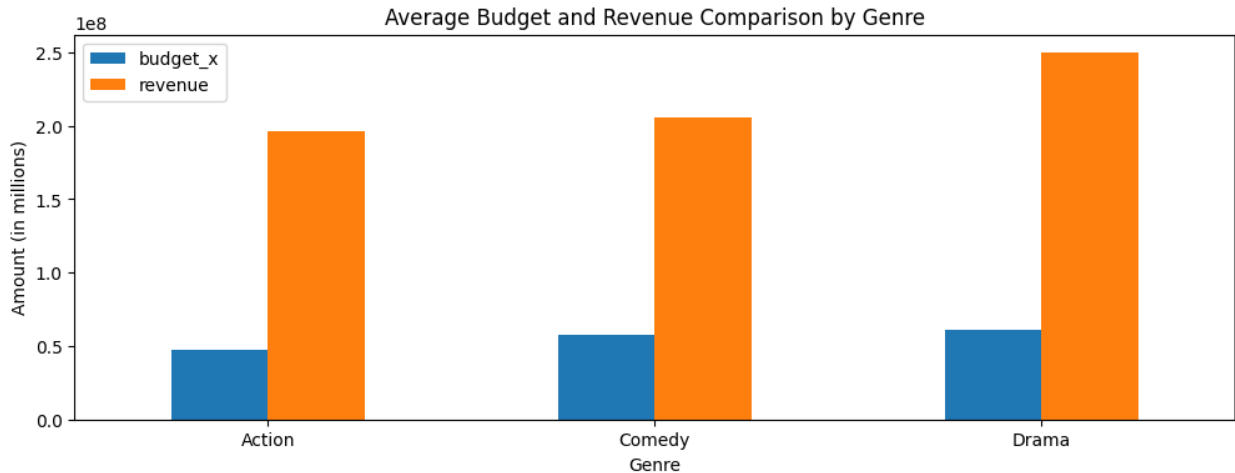
```
# Group by 'genre' and calculate the average budget and revenue
genre_comparison = df_selected_genres.groupby('genre')[['budget_x', 'revenue']].mean()
```

```
# Check the result
print(genre_comparison)
```

	budget_x	revenue
genre		
Action	4.698863e+07	1.962566e+08
Comedy	5.767277e+07	2.052982e+08
Drama	6.068989e+07	2.495831e+08

```
# Plot the comparison of average budget and revenue per genre
genre_comparison.plot(kind='bar', figsize=(12,4 ))
plt.title("Average Budget and Revenue Comparison by Genre")
plt.xlabel("Genre")
plt.ylabel("Amount (in millions)")
plt.xticks(rotation=0)
plt.show()
```

```
# Description
print("The Average Budget and revenue comparison by genre shows that:")
print("- It shows that Drama Genre needs high budget.")
print("- It also shows Action takes less budget than comdey and drama genre")
```



The Average Budget and revenue comparison by genre shows that:

- It shows that Drama Genre needs high budget.
- It also shows Action takes less budget than comedy and drama genre

Year and Trend Analysis


Task: Analyze trends over time

```
# How has the average movie rating changed over the years? Plot the
average rating for each year.

# Group by 'year' and calculate the average rating ('score') for each
year
average_rating_per_year = df.groupby('year')
['score'].mean().reset_index()

# Showing line plot for average movie changed over years
plt.figure(figsize=(12,3))
sns.lineplot(data=average_rating_per_year, x='year', y='score',
marker='o', color='b')
plt.title("Average Movie Rating Over the Years")
plt.xlabel("Year")
plt.ylabel("Average Rating")
plt.xticks(rotation=45)
plt.show()

# Description
print("The Average Movie Rating over the years shows that:")
print("- It shows from 1980 to 2020 the average rating Consistent
Movie Quality")
print("- There down fall after 2020 in movies may be because
expectation of audience")
```

The Average Movie Rating over the years shows that:

- It shows from 1980 to 2020 the average rating Consistent Movie Quality
- There down fall after 2020 in movies may be because expectation of audience

*# Which years had the highest and lowest number of movie releases?
Plot the number of movies released each year*


Group by 'year' and count the number of movies released each year
gb = df.groupby('year').agg({'names': 'count'})

Plot the number of movies released each year using a bar plot

```
plt.figure(figsize=(14,4))
sns.barplot(x=gb.index, y=gb['names'], hue=gb4.index,
palette='viridis')
plt.ylabel('Count of Movies Released')
plt.xlabel('Year')
plt.title('Movies Released Each Year')
plt.xticks(rotation=90, fontsize=7)
plt.show()
```

Find the year with the highest number of movie releases
highest_releases_year = gb['names'].idxmax() *# This gives the year with the highest releases*
print(f"Year with the highest movie releases: {highest_releases_year}")

Find the year with the lowest number of movie releases
lowest_releases_year = gb['names'].idxmin() *# This gives the year with the lowest releases*
print(f"Year with the lowest movie releases: {lowest_releases_year}")



Year with the highest movie releases: 2022

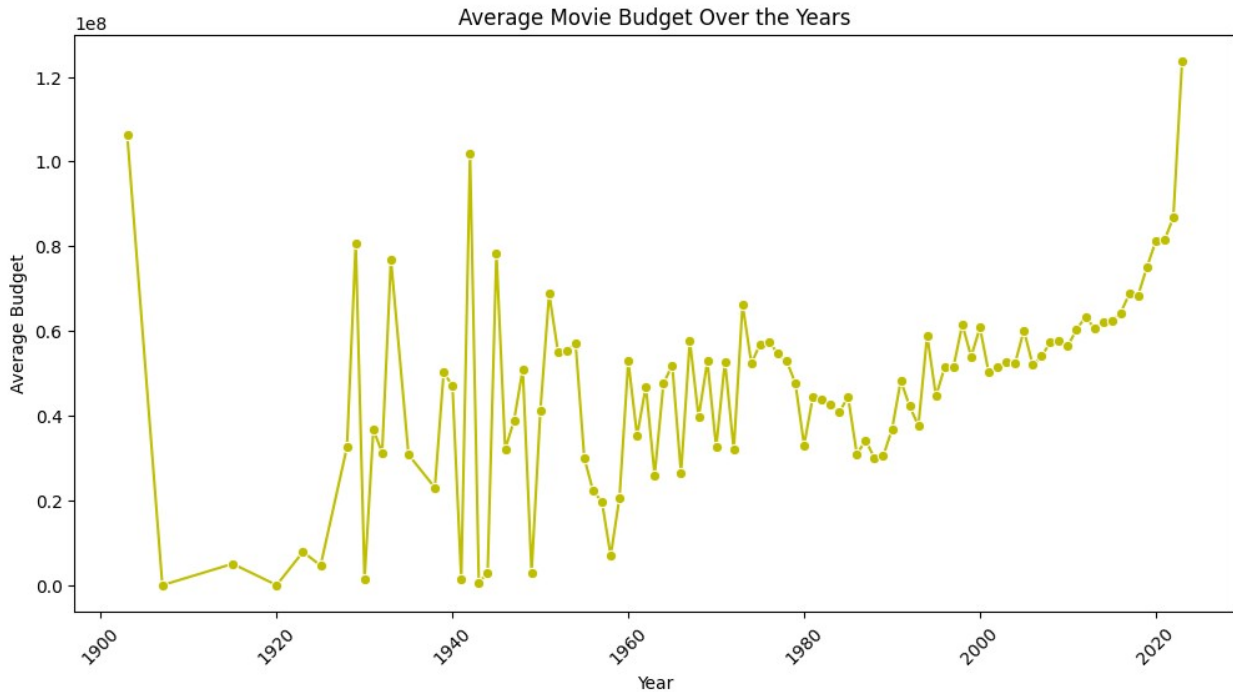
Year with the lowest movie releases: 1903

Do certain years show a higher average budget? Analyze the average budget by year and observe any trends.

Group by 'year' and calculate the average budget for each year
`average_budget_per_year = df.groupby('year')['budget_x'].mean()`

Plot the average budget per year using a line plot

```
plt.figure(figsize=(12, 6))
sns.lineplot(x=average_budget_per_year.index,
y=average_budget_per_year.values, marker='o', color='y')
plt.title("Average Movie Budget Over the Years")
plt.xlabel("Year")
plt.ylabel("Average Budget")
plt.xticks(rotation=45)
plt.show()
```



Multivariate Analysis: Analyze multiple variables together


Combine insights from multiple columns to explore complex relationships

```
# Which genres are most popular in each decade? Create a bar plot
# showing the most frequent genres by decade.

# Extract the decade from the 'year' column
df['decade'] = (df['year'] // 10) * 10 # Dividing year by 10 and
# multiplying by 10 to get the start of the decade

# Plot the number of movies released each decade
gb = df.groupby('decade').agg({'genre': 'count'})
plt.figure(figsize = (12, 4))
sns.barplot(x = gb.index, y = gb['genre'], data = gb, hue = gb.index,
palette = 'viridis')
plt.ylabel('Count of Movies Released')
plt.xlabel('Year')
plt.title('Movies Released each Decade')
plt.xticks(rotation = 90, fontsize = 7)
plt.show()

# Description
print("The Movies released each decade plot shows that:")
print("- It shows in 2010 most movies was popular ")
```



The Movies released each decade plot shows that:

- It shows in 2010 most movies was popular

Analyze the influence of `budget_x`, `genre`, and `country` on revenues using a heatmap.

Select the relevant columns and drop any missing values

```
df_corr = df[['budget_x', 'genre', 'country', 'revenue']].dropna()
```

Convert categorical columns into numeric values

```
df_corr['genre'] = pd.factorize(df_corr['genre'])[0]
```

```
df_corr['country'] = pd.factorize(df_corr['country'])[0]
```

Calculate the correlation between the variables

```
corr_matrix = df_corr.corr()
```

```
print(corr_matrix)
```

	budget_x	genre	country	revenue
budget_x	1.000000	-0.008550	0.159005	0.680372
genre	-0.008550	1.000000	-0.018373	-0.010409
country	0.159005	-0.018373	1.000000	0.145098
revenue	0.680372	-0.010409	0.145098	1.000000

Plot the heatmap

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
```

```
plt.title("Correlation Heatmap: Budget, Genre, Country, and Revenue")
```

```
plt.show()
```


Description

```
print("The correlation Heatmap shows that:")
```

```
print("- Budget and revenue are positively correlated (0.67),  
suggesting that higher-budget movies tend to generate more revenue")
```

```
print("- Genre and country have very weak correlations with other
```

```
variables, indicating that the genre or the country of production does
not strongly impact budget or revenue.")
print("-Country has a weak positive correlation with revenue
suggesting that movies from certain countries might tend to have
slightly higher revenue.")
```



The correlation Heatmap shows that:

- Budget and revenue are positively correlated (0.67), suggesting that higher-budget movies tend to generate more revenue
- Genre and country have very weak correlations with other variables, indicating that the genre or the country of production does not strongly impact budget or revenue.
- Country has a weak positive correlation with revenue suggesting that movies from certain countries might tend to have slightly higher revenue.

Additional Questions Based on Dataset

```
# Which countries produce the highest-rated movies on average?

# Group by 'country' and calculate the average rating ('score') for
each country
average_rating_by_country = df.groupby('country')
['score'].mean().sort_values(ascending=False).reset_index()


# Display the top 10 countries with the highest average ratings
top_countries = average_rating_by_country.head(10)
print("Top 10 Countries with Highest Average Movie Ratings:")
print(top_countries)
```

Top 10 Countries with Highest Average Movie Ratings:

	country	score
0	SU	79.800000
1	PR	76.000000
2	XC	76.000000
3	CZ	72.500000
4	DO	72.000000
5	PY	71.000000
6	MU	70.000000
7	MX	69.771429
8	GT	68.000000
9	DK	67.333333

```
# Plot the results using a barplot
plt.figure(figsize=(12, 6))
sns.barplot(data=top_countries, x='country',
y='score', hue=top_countries.index, palette='viridis')
plt.title("Top 10 Countries with Highest Average Movie Ratings")
plt.xlabel("Country")
plt.ylabel("Average Rating")
plt.xticks(rotation=90)
plt.show()
```

```
# Description
print("The Highest Average movie ratings shows that:")
print("- It shows that SU has highest average movie ratings")
```



The Highest Average movie ratings shows that:

- It shows that SU has highest average movie ratings

Does the original language ('orig_lang') correlate with ratings?

Group by 'orig_lang' and calculate the average rating for each language

```
avg_rating_by_lang = df.groupby('orig_lang')
['score'].mean().sort_values(ascending=False).reset_index()
```

```
top= avg_rating_by_lang.head(20)
```

Display the results

```
print("Average Ratings by Original Language:")
```

```
print(top)
```

Average Ratings by Original Language:

	orig_lang	score
0	Irish	76.000000
1	Czech	75.500000
2	Kannada	75.500000
3	Galician	75.000000
4	Dzongkha	74.000000
5	Gujarati	73.000000
6	Romanian	72.500000
7	Ukrainian	71.500000
8	Arabic	71.500000
9	Latin	71.000000
10	Catalan, Valencian	71.000000

```


11         Danish 69.304348
12         Persian 68.200000
13         Hindi 67.653846
14         Indonesian 67.636364
15     Bokmål, Norwegian, Norwegian Bokmål 67.000000
16         Japanese 66.899160
17         Macedonian 66.000000
18         Spanish, Castilian 65.188917
19         Swedish 65.000000

```

```

# Plot the results using a barplot
plt.figure(figsize=(12,4))
sns.barplot(data=top, x='orig_lang', y='score', hue = top.index,
palette='viridis')
plt.title("Average Movie Ratings by Original Language")
plt.xlabel("Original Language")
plt.ylabel("Average Rating")
plt.xticks(rotation=90)
plt.show()

```



```
df["status"]
```

```

0      Released
1      Released

```



```
2      Released
3      Released
4      Released
```

```
...
10173    Released
10174    Released
10175    Released
10176    Released
10177    Released
```

```
Name: status, Length: 10178, dtype: object
```

```
# How does `status` (e.g., released, post-production) influence revenues?
```

```
# Filter out movies with invalid revenue values (e.g., 0 or NaN)
df_filtered_revenue = df[df['revenue'] > 0] # Filter out rows with zero revenue
```

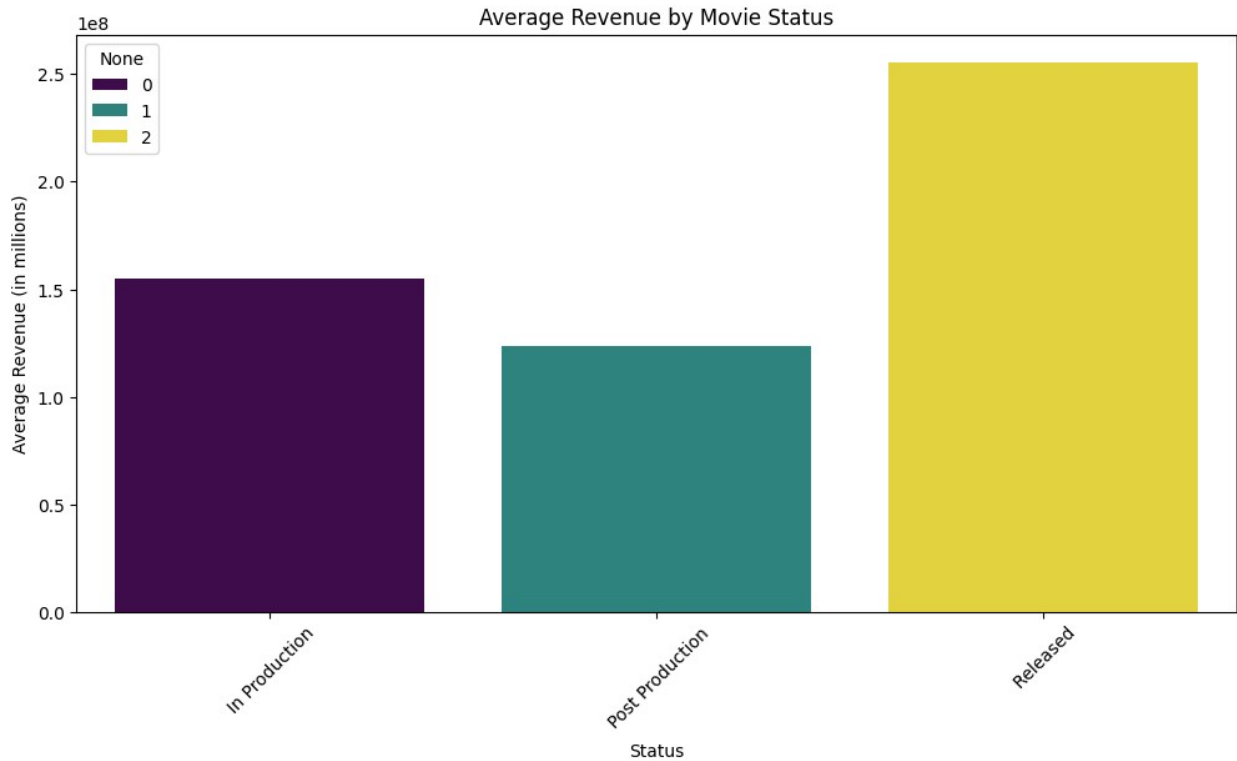
```
# Group by 'status' and calculate the average revenue for each status
avg_revenue_by_status = df_filtered_revenue.groupby('status')
['revenue'].mean().reset_index()
```

```
# Plot the average revenue by status
```

```
plt.figure(figsize=(12, 6))
sns.barplot(data=avg_revenue_by_status, x='status',
y='revenue', hue=avg_revenue_by_status.index, palette='viridis')
plt.title("Average Revenue by Movie Status")
plt.xlabel("Status")
plt.ylabel("Average Revenue (in millions)")
plt.xticks(rotation=45)
plt.show()
```

```
# Description
```

```
print("The Average revenue by movie status shows that:")
print("- Status which is Released have highest avergae revenue ")
```



The Average revenue by movie status shows that:
- Status which is Released have highest average revenue

Insights and Summary

Task: Summarize key findings

1. Budget and Revenue Correlation:

- There is a **strong positive correlation** (0.67) between **budget** and **revenue**, indicating that movies with higher budgets tend to generate more revenue. This suggests that investing more in movie production is likely to lead to higher financial returns. However, it is important to note that while a higher budget increases the chances of higher revenue, it does not guarantee success. Some high-budget films may still fail to generate expected returns, and there are mid-budget films that perform better than anticipated.

2. Consistent Movie Ratings Over Time:

- From **1980 to 2020**, the average **IMDB ratings** have remained relatively **consistent**, showing that movie quality, as measured by user ratings, has not drastically changed over time. However, there has been a **slight decline after 2020**, which could be attributed to shifting audience expectations, a decrease in movie-going experiences due to external factors (such as the pandemic), or a change in the types of content being produced. This trend suggests that while movie quality has remained stable, audience preferences and expectations are evolving.

3. Genre-Specific Trends in Budget and Ratings:

- Certain **genres** like **Fantasy**, **Drama**, and **Crime** tend to receive **higher average ratings**, while **Action** movies often have **lower budgets** compared to **Comedy** and **Drama** films. This shows that more dramatic or fantasy-driven films may receive better audience ratings, possibly due to stronger narratives or unique storytelling. In contrast, Action films often have lower production costs but can still generate significant box-office returns, showing that action-based content might be more accessible or widely popular despite its lower production costs.